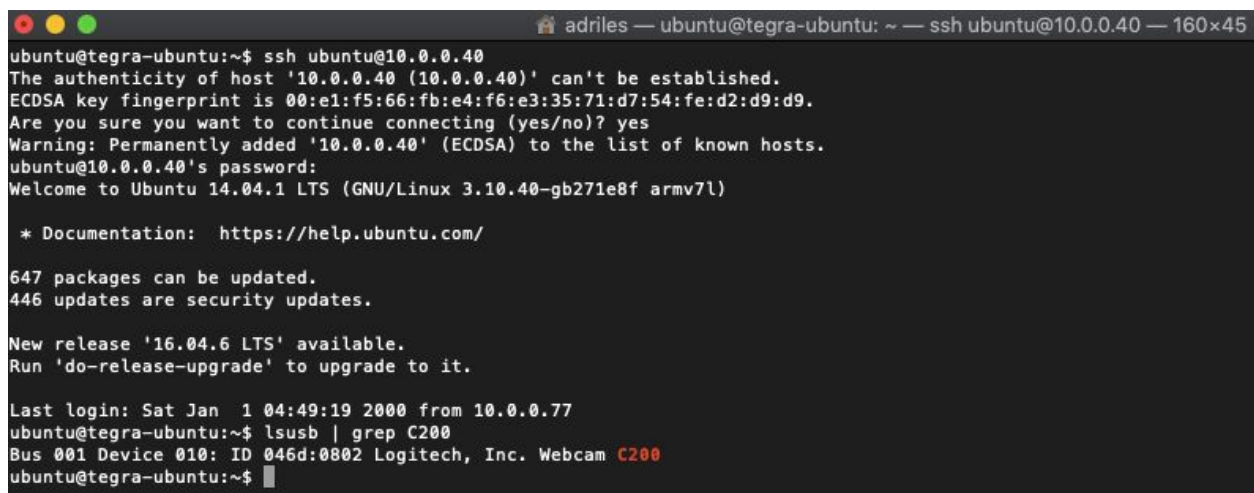


Part 1

Here we simply needed to verify that the Logitech Webcam C200 is able to connect and be recognized by the Jetson TK1. We use the “lsusb” kernel driver configuration tool as seen below to verify that the camera is recognized:

A terminal window titled 'adriels — ubuntu@tegra-ubuntu: ~ — ssh ubuntu@10.0.0.40 — 160x45'. The terminal shows an SSH session from ubuntu@tegra-ubuntu to ubuntu@10.0.0.40. It displays the host's ECDSA key fingerprint, a warning about adding it to the list of known hosts, and the user's password. After logging in, it shows the Ubuntu 14.04.1 LTS welcome message, documentation link, package update information (647 packages, 446 security updates), and a new release '16.04.6 LTS' available. The user then runs 'lsusb | grep C200', which returns 'Bus 001 Device 010: ID 046d:0802 Logitech, Inc. Webcam C200'.

```
ubuntu@tegra-ubuntu:~$ ssh ubuntu@10.0.0.40
The authenticity of host '10.0.0.40 (10.0.0.40)' can't be established.
ECDSA key fingerprint is 00:e1:f5:66:fb:e4:f6:e3:35:71:d7:54:fe:d2:d9:d9.
Are you sure you want to continue connecting (yes/no)? yes
Warning: Permanently added '10.0.0.40' (ECDSA) to the list of known hosts.
ubuntu@10.0.0.40's password:
Welcome to Ubuntu 14.04.1 LTS (GNU/Linux 3.10.40-gb271e8f armv7l)

 * Documentation:  https://help.ubuntu.com/

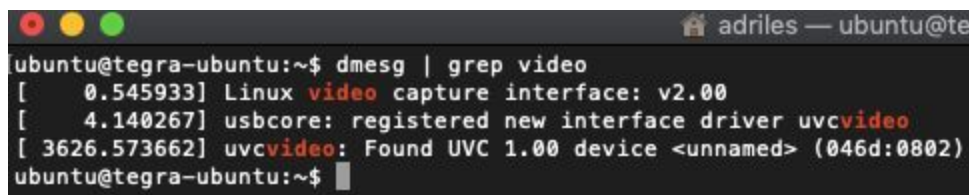
647 packages can be updated.
446 updates are security updates.

New release '16.04.6 LTS' available.
Run 'do-release-upgrade' to upgrade to it.

Last login: Sat Jan  1 04:49:19 2000 from 10.0.0.77
ubuntu@tegra-ubuntu:~$ lsusb | grep C200
Bus 001 Device 010: ID 046d:0802 Logitech, Inc. Webcam C200
ubuntu@tegra-ubuntu:~$
```

Figure 1: Showing that the Webcam C200 device is recognized as a connected device to the USB bus.

We then also use the “dmesg” command to see that the UVC driver is loaded:

A terminal window titled 'adriels — ubuntu@te' showing the command 'dmesg | grep video'. The output shows three lines of kernel messages: '[0.545933] Linux video capture interface: v2.00', '[4.140267] usbcore: registered new interface driver uvcvideo', and '[3626.573662] uvcvideo: Found UVC 1.00 device <unnamed> (046d:0802)'.

```
ubuntu@tegra-ubuntu:~$ dmesg | grep video
[ 0.545933] Linux video capture interface: v2.00
[ 4.140267] usbcore: registered new interface driver uvcvideo
[ 3626.573662] uvcvideo: Found UVC 1.00 device <unnamed> (046d:0802)
ubuntu@tegra-ubuntu:~$
```

Figure 2: Showing that the kernel recognized a new interface driver “uvcvideo” and registered it with the usbcore. Furthermore, the “uvcvideo” driver found a UVC 1.00 device, which is not explicitly said to be our Webcam C200, but implicitly must be.

Part 2

Here we needed to install either Camorama or Cheese and do a simple snapshot capture. This proved to us that the camera had at least its basic functionality. Camorama and Cheese were already installed on this Jetson, but I did have to uninstall the gnome desktop and install the gtk desktop to make Camorama functional. I did not use Cheese once Camorama was working.

Screenshots of Camorama, as well as a few images in which I played with the color balance are shown below:

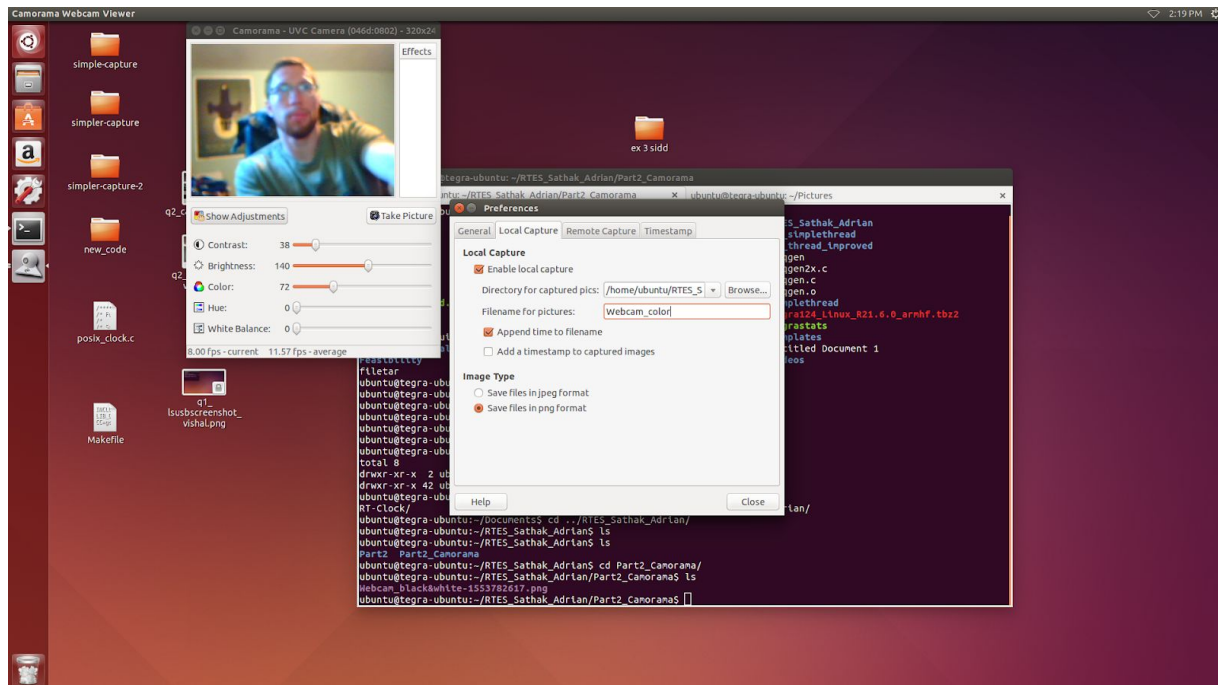


Figure 3: Screenshot of Camorama installed and functional on the Jetson TK1

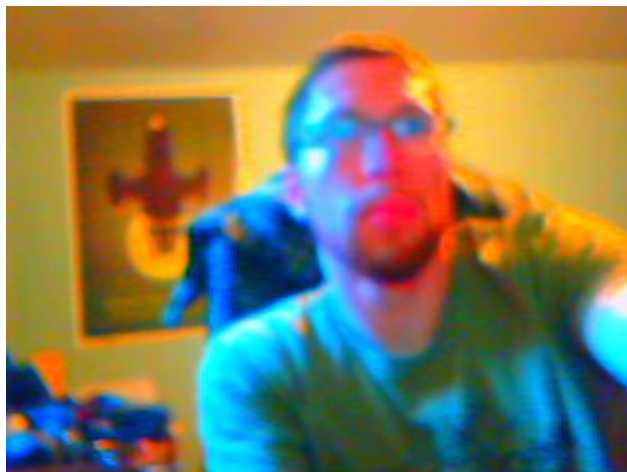


Figure 4: Snapshot with Camorama, turning up the color balance to maximum

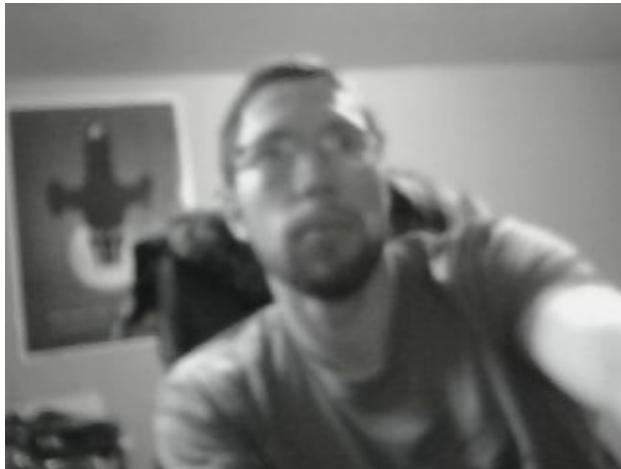


Figure 5: Snapshot with Camorama, turning down the color balance to minimum

Part 3

- Using your verified Logitech C200 camera on a DE1-SoC, Raspberry Pi or Jetson, verify that it can stream continuously using to a raw image buffer for transformation and processing using example code such as [simple-capture](#), [simpler-capture](#), or [simpler-capture-2](#). Provide a screenshot to prove that you got continuous capture to work.

Compilation of simple-capture was straightforward since it doesn't require OpenCV. For screenshots of compiling simple-capture and simpler-capture, look in the Part3 folder of the Exercise 4 submission.

simple-capture

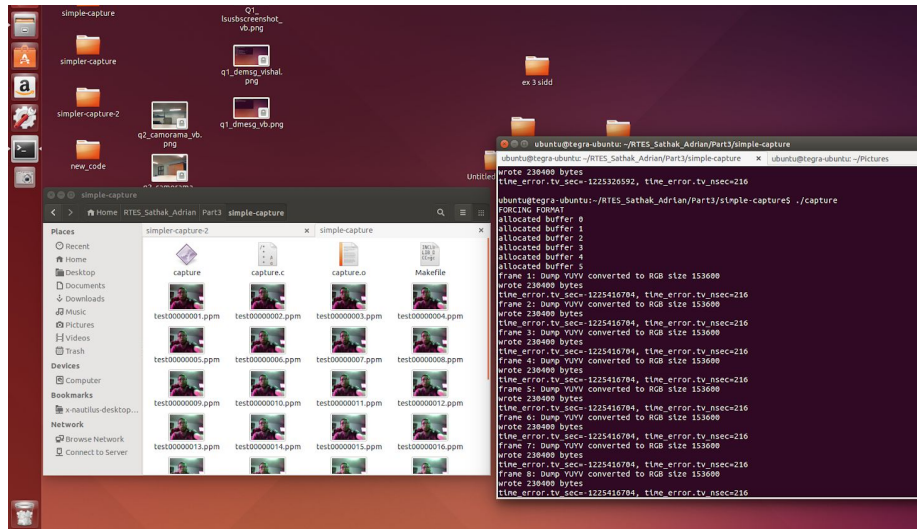


Figure 6: Running simple-capture, a stream of images written to the disk

The output stream of images from simple-capture are included in the submission folder at: Part3/streaming_test1:simple-capture. The output files from the streaming capture are .ppm files.

simpler-capture

Using simpler-capture required setup of OpenCV, which started a good deal of forum surfing and ultimately was a bit frustrating. Initially I tried updating to the most recent version of OpenCV, currently at 4.0. I was unable to compile simpler-capture however, because the package -libopencv-nonfree couldn't be found. I tried recompiling OpenCV but was met with

errors since Ubuntu 14.04 LTS has version 2.8 of cmake installed, and the recompilation required at least version 3.5.

I purged cmake from the system and installed cmake 3.8, but even with this version, I wasn't able to recompile. Cmake had installed to /usr/local/bin, and the build process was only looking in /usr/bin for cmake source. Adding symbolic links didn't solve this build issue. Eventually after speaking with the TAs, I learned that OpenCV versions 3 and 4 were problematic, and that I should downgrade to version 2.

I purged OpenCV and cmake, reinstalled cmake 2.8 and then downloaded, built, and installed OpenCV2.4.13. On my first try, simpler-capture built as shown here:

Running simpler-capture is shown below:

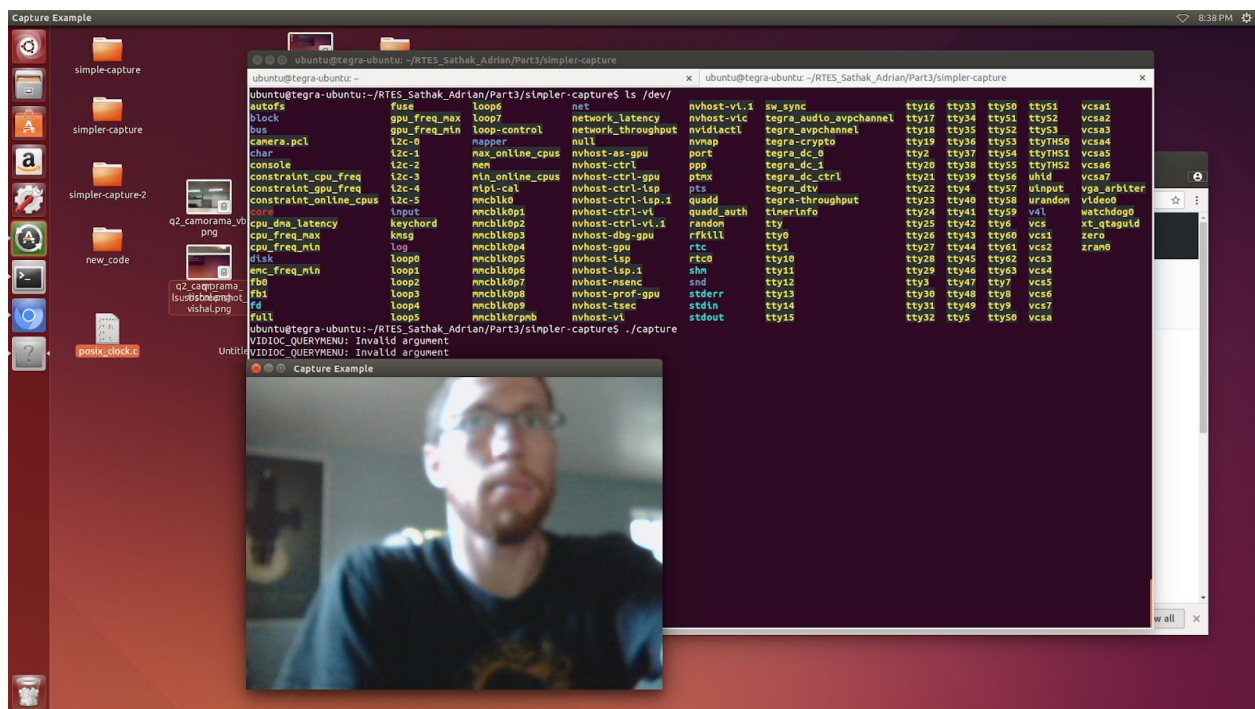


Figure 7: Running simpler-capture

Part 4

- Choose a continuous transformation OpenCV example from [computer-vision](#) such as the [canny-interactive](#), [hough-interactive](#), [hough-elliptical-interactive](#), or [stereo-transform-improved](#) or the from the same 4 transforms in [computer vision cv3 tested](#). Show a screenshot to prove you built and ran the code. Provide a detailed explanation of the code and research uses for the continuous transformation by looking up API functions in the OpenCV manual (<http://docs.opencv.org>)

I compiled and ran the simple-canny-interactive, simple-hough-interactive, and simple-hough-elliptical-interactive programs. See the “compiling” screenshots in the Part4 folder of the submission.

The simple-hough-interactive program takes an image, finds the straight lines in that image, and then overlays markers of those straight lines into the display. Explanation of the simple-hough-interactive code upon entering main() follows.

Before the while(1) loop:

Firstly, a window is created for the camera's image will be displayed. Parameters are given such that the window is named “Capture Example” and the window is auto-sized to fit the size of the image. The dimensions that will be used for this are given as #defines, 640 x 480 pixels.

Instances of two classes (CvCapture and IplImage) are declared. The CvCapture class is used to facilitate capture of video from a camera. The IplImage class is from the Intel Image Processing Library, and is used to process a single frame. An int is created to identify the device as video”0”, and then a few Mat (matrix) objects are declared. Finally a vector of type Vec4i is defined.

After command-line arguments are processed, the capture is initialized with a call to cvCreateCameraCapture() with the device identifier passed as the parameter. Then the height and width of the capture are set to the #defined dimensions for height and width resolutions with calls to cvSetCaptureProperty() and the #defines mentioned before.

Then begins the while(1) loop:

A frame (of type IplImage) is captured with a call to cvQueryFrame() with the CvCapture class object passed as a parameter. This call gets the last frame from the CvCapture classes internal buffer, that's been filling since the cvCreateCameraCapture() call.

The frame is added to a new matrix “mat_frame”. Then this matrix has an edge-detection function applied to it called Canny(). This takes the matrix which contains the last frame

“mat_frame” and applies the transformation within the threshold limits [50, 200], outputting the result to the matrix canny_frame. This transformation is an edge detector, (this transformation alone can be seen in the simple-canny-interactive program) and would display edges as white lines on a black background.

This “canny_frame” matrix then passed to another transformation function: HoughLines(). This function takes the edge lines detected in canny_frame and tries to detect the straight lines among those edges. The call to HoughLines() in simple-hough-interactive uses 1 pixel and 1 degree for resolution, 50 intersections as the threshold for detection, and 50 and 10 for the distance resolutions of rho and theta respectively.

The resulting straight lines detected are put into a vector Vec4i. In a for-loop, each of these lines is drawn onto the matrix “mat_frame”, using the point vertices the color red (given as a RGB scalar), and a thickness of 3 pixels. Furthermore the lines are drawn as an antialiased line, which are specifically rendered for smoothness.

After the lines have been added to the matrix, and if the frame wasn't null to begin with, the frame is shown to the “Capture Example” window. A cvWaitKey() call then waits for 10 milliseconds. If a key is detected and that key is “escape” (key 27), the while(1) loop ends and the program exits. Otherwise the while(1) loop continues, more frames are captured and transformed first to find edges and then straight lines, and finally displayed to the “Capture Example” window.

Upon detection of the “escape” key, the capture is ended with a call to cvReleaseCapture() and resources are freed. Then the “Capture Example” display window is closed with a call to cvDestroyWindow(), and the program exits.

simple-hough-interactive uses

The simple-hough-interactive continuous transform has many possible uses, particularly in research. For instance, airborne imagery of a forest with many Pine trees can be clouded with all the branches and foliage, but if you just want to find the tree trunks the HoughLine() transform performed in simple-hough-interactive would be a good place to start. Other types of images in which you might want to find straight lines could be satellite data in order to map roads for cartographic purposes, or sidescan sonar data in which you could be looking for wrecked ships and planes.

Screenshots of each of the transformations (simple-canny-interactive, simple-hough-interactive, and simple-hough-elliptical-interactive) follow. Additionally, note that I took screenshots with an open terminal window running the “top” command. This allows us to see the CPU usage of each program as they run.

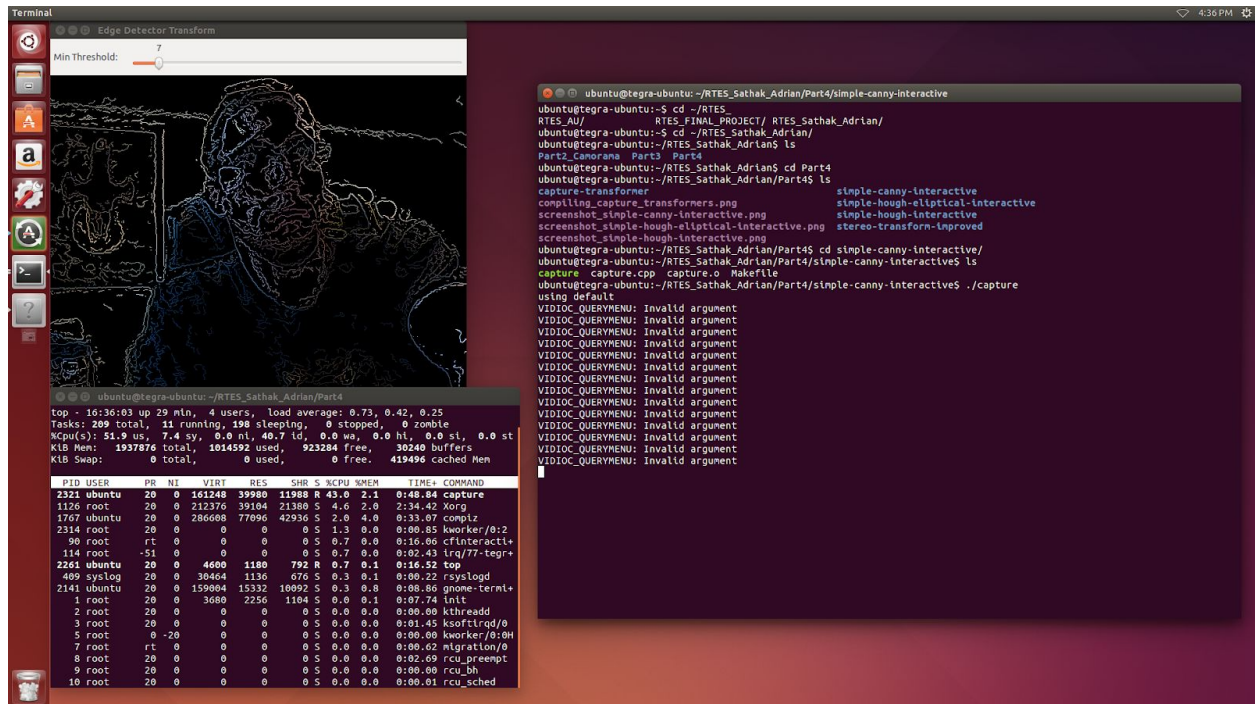


Figure 8: Running simple-canny-interactive

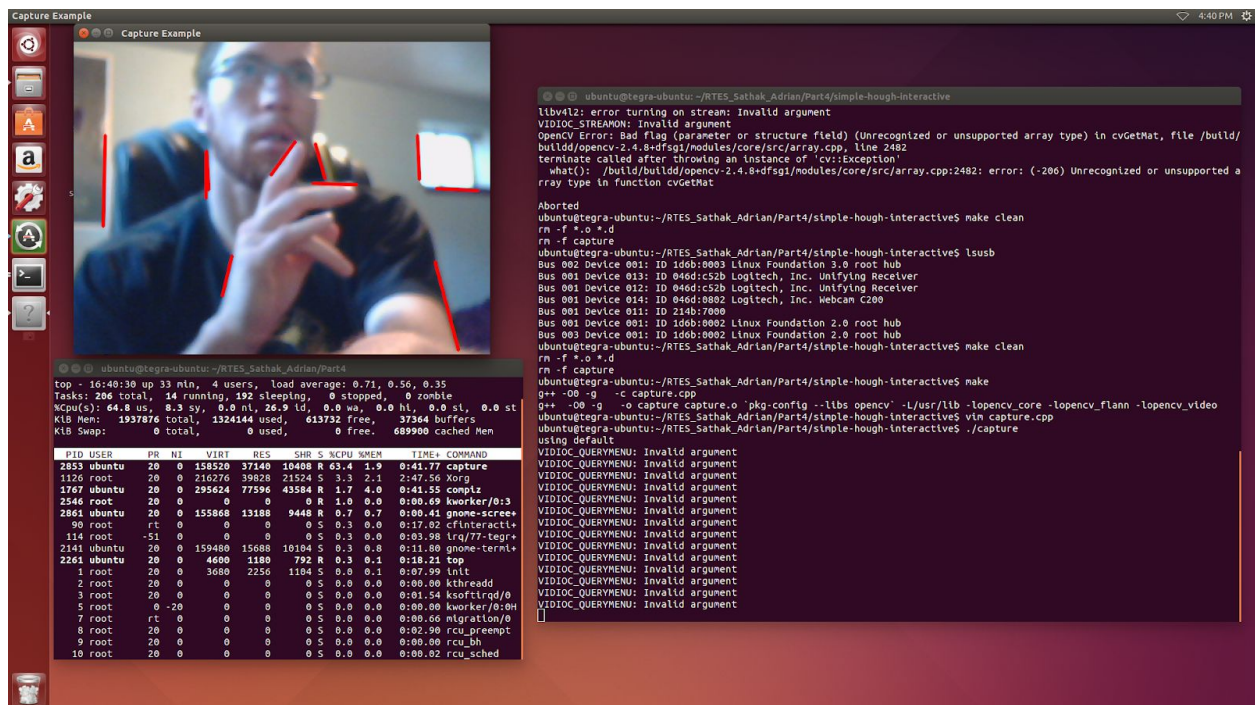


Figure 9: Running simple-hough-interactive

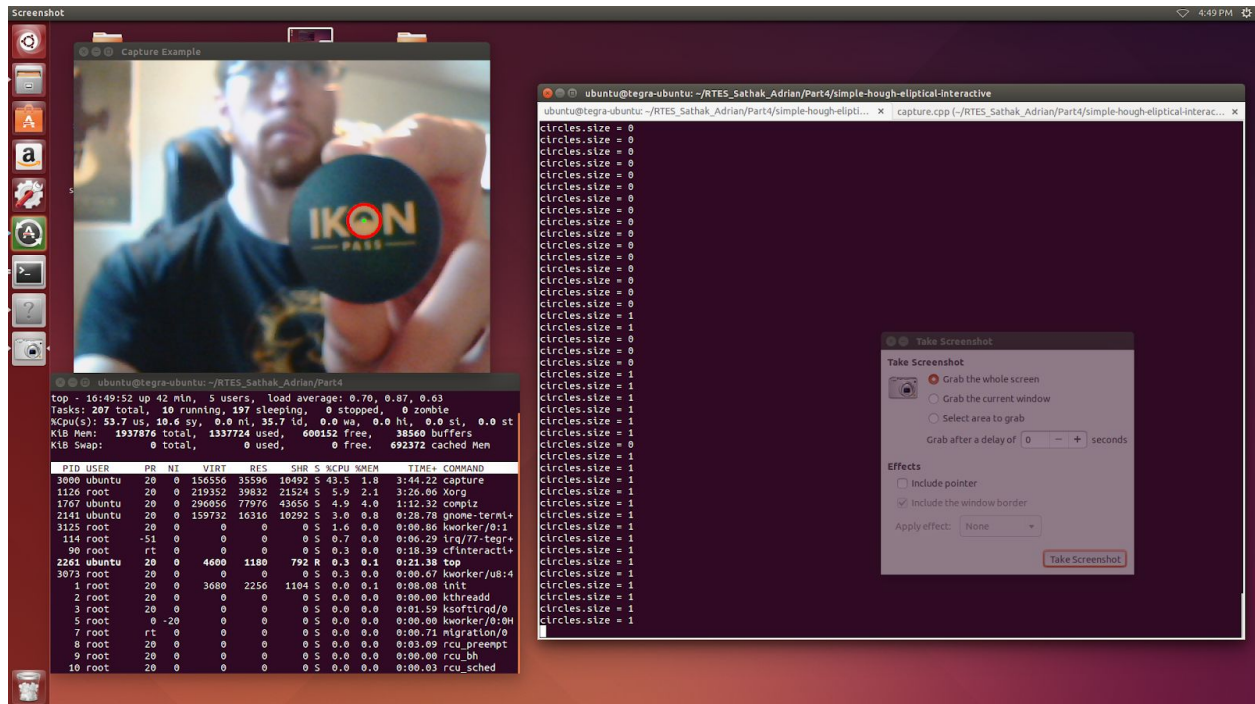


Figure 10: Running simple-hough-elliptical-interactive

- for stereo-transform-improved either implement or explain how you could make this work continuously rather than snapshot only.

After downloading the code from Prof. Siewert's website, I was a bit confused by the task we were assigned. There are two programs included in the directory, `capture.cpp` and `capture_stereo.cpp`. The `capture.cpp` program already has a continuous stream capability, and the `capture_stereo.cpp` program also encapsulates the abilities to make continuous Canny (edge detection) and Hough (straight line) transformations within `while(1)` loops. What doesn't exist is solely the ability to save images to the disc with these applied transformations.

It should be noted that files are only saved in `capture_stereo.cpp` when the "escape" key is pressed, in a snapshot fashion. I have doubts this was the intended update to the program, but images can be continuously written to the disc by removing the "if" condition that looks for "`char c == ESC_KEY`". See `capture_stereo.cpp` in the stereo-transform-improved folder in the Part4 folder of the submission.

Part 5

Using a Logitech C200, choose 3 real-time interactive transformations to compare in terms of average frame rate at a given resolution for a range of at least 3 resolutions in a common aspect ratio (e.g. 4:3 for 1280x960, 640x480, 320x240, 160x120, 80x60) by adding time-stamps to analyze the potential throughput. Based on average analysis, pick a reasonable soft real-time deadline (e.g. if average frame rate is 12 Hz, choose a deadline of 100 milliseconds to provide some margin) and convert the processing to SCHED_FIFO and determine if you can meet deadlines with predictability and measure jitter in the frame rate relative to your deadline.

Ans:

Three threads have been created, with each thread running its own version of a transform, simple canny, hough and elliptical hough respectively.

The code has been designed such that three semaphores are used, one for each thread. Each thread has two while loops, one inner and one outer. The outer loop is a while(1), and the inner is a while loop for 50 frame counts.

The semaphore for thread 1(simple canny) has been initialized with value 1, and the other two semaphores with value 0. Each thread has a 'sem_wait' at the beginning, and releases the semaphore for the next thread at the end of its inner while loop. The chain continues so on and so forth allowing a type of FIFO scheduling using three semaphores' logic.

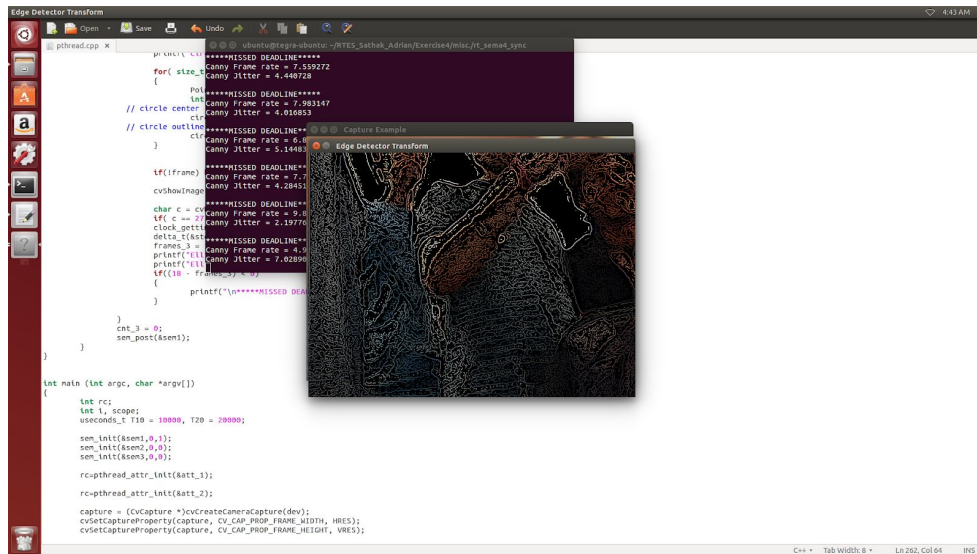
All screenshots below have been listed for 640 x 480 resolution.

The table below lists the FPS set for each transform per resolution:

Transform	640 x 480	320 x 240	160 x 120
Canny	10 fps	20 fps	25 fps
Hough	13 fps	32 fps	33 fps
Hough Elliptical	18 fps	20 fps	20 fps

The table below lists the deadline set for each transform per resolution:

Transform	640 x 480	320 x 240	160 x 120
Canny	125 ms	56 ms	45 ms
Hough	100 ms	33 ms	33 ms
Hough Elliptical	67 ms	56 ms	56 ms



In the image above, one can see the canny transform of an image, with the output in the background, and a few deadlines missed.

As we can see in the screenshots below, some deadlines are missed occasionally.

```

*****MISSED DEADLINE*****
Canny Frame rate = 13.436346
Canny Jitter = 11.563654
Canny Frame rate = 22.307196
Canny Jitter = 2.692804
Canny Frame rate = 26.630802
Canny Jitter = -1.630802

*****MISSED DEADLINE*****
Canny Frame rate = 24.742117
Canny Jitter = 0.257883
Canny Frame rate = 26.613733
Canny Jitter = -1.613733

*****MISSED DEADLINE*****
Canny Frame rate = 26.692232
Canny Jitter = -1.692232

*****MISSED DEADLINE*****
Canny Frame rate = 26.504812
Canny Jitter = -1.504812

*****MISSED DEADLINE*****
Canny Frame rate = 25.289083
Canny Jitter = -0.289083

*****MISSED DEADLINE*****
Canny Frame rate = 24.396296
Canny Jitter = 0.603704
Canny Frame rate = 25.221106
Canny Jitter = -0.221106

*****MISSED DEADLINE*****
Canny Frame rate = 25.459385
Canny Jitter = -0.459385

*****MISSED DEADLINE*****
Canny Frame rate = 25.663509
Canny Jitter = -0.663509

*****MISSED DEADLINE*****
Canny Frame rate = 24.766422
Canny Jitter = 0.233578
Canny Frame rate = 21.742636
Canny Jitter = 3.257364
Canny Frame rate = 25.379530
Canny Jitter = -0.379530

```


The average frame rate for each resolution and transform has been observed, and an arbitrarily ranged value has been selected for calculating the deadline. If a deadline is missed, a "MISSED DEADLINE" tag is printed, and the relative jitter is calculated and displayed.

The final relative jitter can be predicted to a particular extent, and it is partly depends on the code as well the other processes running on the board at the time of code execution.

It was observed that even if a short gap is kept in between executing the code consecutively, the jitter does vary a bit, leading me to believe that the board has a number of other subsidiary processes running as well.